

Assignment14: LSTM

LSTM model on Amazon Fine Food Reviews

Objective:

- Apply LSTM model with different architectures & number of layers on Amazon food reviews.
- Plots Training and testing error plot.

```
In [0]: import warnings
warnings.filterwarnings("ignore")
```

```
In [0]: %matplotlib inline
import re
from datetime import datetime
import math
import pandas as pd
import numpy as np
import nltk
import string
import pickle
import matplotlib.pyplot as plt

import seaborn as sns

from sklearn.feature_extraction.text import CountVectorizer
from nltk.stem.porter import PorterStemmer

from nltk.corpus import stopwords
from nltk.stem import PorterStemmer
```

```
from nltk.stem.wordnet import WordNetLemmatizer

from gensim.models import Word2Vec
from gensim.models import KeyedVectors
```

```
In [0]: # Importing libraries
from keras.models import Sequential
from keras.layers import Dense
from keras.layers import LSTM
from keras.layers.embeddings import Embedding
from keras.preprocessing import sequence
from keras.layers import Dropout
# fix random seed for reproducibility
np.random.seed(7)
```

Using TensorFlow backend.

```
In [0]: %matplotlib inline
import matplotlib.pyplot as plt
import numpy as np
import time
# https://gist.github.com/greydanus/f6eee59eaf1d90fcb3b534a25362cea4
# https://stackoverflow.com/a/14434334
# this function is used to update the plots for each epoch and error
def plt_dynamic(x, vy, ty):
    fig = plt.figure( facecolor='c', edgecolor='k')
    plt.plot(x, vy, 'b', label="Validation Loss")
    plt.plot(x, ty, 'r', label="Train Loss")
    plt.xlabel('Epochs')
    plt.ylabel('Categorical Crossentropy Loss')
    plt.legend()
    plt.grid()
    plt.show()
```

To upload the reviews.zip file from system to google colab

```
In [0]: # Reference link
# https://towardsdatascience.com/3-ways-to-load-csv-files-into-colab-7c14fcbdc92
# Code to read csv file into Colaboratory:
!pip install -U -q PyDrive
from pydrive.auth import GoogleAuth
from pydrive.drive import GoogleDrive
from google.colab import auth
from oauth2client.client import GoogleCredentials
# Authenticate and create the PyDrive client.
auth.authenticate_user()
gauth = GoogleAuth()
gauth.credentials = GoogleCredentials.get_application_default()
drive = GoogleDrive(gauth)
```

```
100% |████████████████████████████████████████████████████████████████████████████████| 993kB 19.7MB/s
Building wheel for PyDrive (setup.py) ... done
```

```
In [0]: link = 'https://drive.google.com/file/d/1p5YpkUos0Rd6Hsae1qaH0VsPtnLdlnc0/view?usp=sharing'
```

```
In [0]: fluff, id = link.split('=')
id='1p5YpkUos0Rd6Hsae1qaH0VsPtnLdlnc0'
print (id)
```

```
1p5YpkUos0Rd6Hsae1qaH0VsPtnLdlnc0
```

```
In [0]: downloaded = drive.CreateFile({'id':id})
downloaded.GetContentFile('Reviews.csv')
```

```
In [0]: amz = pd.read_csv('Reviews.csv')
# dimensions of dataset and columns name
print(amz.shape)
#print(amz1.shape)
print(amz.columns)
```

```
(568454, 10)
Index(['Id', 'ProductId', 'UserId', 'ProfileName', 'HelpfulnessNumerato
```

```
r',  
      'HelpfulnessDenominator', 'Score', 'Time', 'Summary', 'Text'],  
      dtype='object')
```

Given reviews with Score>3 a positive rating, and reviews with a score<3 a negative rating. If score is equal to 3, it is considered as neutral score.

```
In [0]: # Processing  
#Give reviews with Score>3 a positive rating, and reviews with a score<  
3 a negative rating.  
  
def score_part(x):  
    if x < 3:  
        return 0  
    return 1  
  
actualScore = amz['Score']  
#print(actualScore)  
New_score = actualScore.map(score_part)  
#print(New_score)  
amz['Score']=New_score
```

```
In [0]: print(amz.shape)  
amz.head(2)
```

```
(568454, 10)
```

Out[0]:

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	Helpfulnes
0	1	B001E4KFG0	A3SGXH7AUHU8GW	delmartian	1	1

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	Helpfulnes
1	2	B00813GRG4	A1D87F6ZCVE5NK	dll pa	0	0

Data Pre-processing on raw data:

Every datasets contains some unwanted data.Raw data is preprocessed by removing duplication.

```
In [0]: #Processing of ProductId
#Sorting data according to ProductId in ascending order
sorted_data=amz.sort_values('ProductId', axis=0, ascending=True,
                           inplace=False, kind='quicksort',
                           na_position='last')
#sorted_data.head() # printing sorted data
# To check the duplications in raw data
dupli=sorted_data[sorted_data.duplicated(["UserId","ProfileName","Time"
,"Text"])]
print(dupli.head(2))
# Remove Deduplication of entries
final=sorted_data.drop_duplicates(subset={"UserId","ProfileName","Time"
,"Text"}, keep='first', inplace=False)
final.shape
#Checking to see how much % of data still remains
(final['Id'].size*1.0)/(amz['Id'].size*1.0)*100
final=final[final.HelpfulnessNumerator<=final.HelpfulnessDenominator]
#Before starting the next phase of preprocessing lets see the number of
entries left
print(final.shape)
```

	Id	ProductId	UserId	\
171222	171223	7310172001	AJD41FBJD9010	
171153	171154	7310172001	AJD41FBJD9010	

		ProfileName	HelpfulnessNumera
tor	\		
171222	N. Ferguson	"Two, Daisy, Hannah, and Kitten"	
	1		
171153	N. Ferguson	"Two, Daisy, Hannah, and Kitten"	
	0		

	HelpfulnessDenominator	Score	Time	\
171222	1	1	1233360000	
171153	0	1	1233360000	

		Summary	\
171222	best dog treat-- great for training---	all do...	
171153	best dog treat-- great for training---	all do...	

		Text
171222	Freeze dried liver has a hypnotic effect on do...	
171153	Freeze dried liver has a hypnotic effect on do...	

(393931, 10)

```
In [0]: final = final[final['ProductId'] != '7310172001']
```

Text Preprocessing:

```
In [0]: import nltk
nltk.download('stopwords')

[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data] Unzipping corpora/stopwords.zip.
```

```
Out[0]: True
```

```
In [0]:
```

```

stop = set(stopwords.words('english')) #set of stopwords
sno = nltk.stem.SnowballStemmer('english') #initialising the snowball s
temmer

def cleanhtml(sentence): #function to clean the word of any html-tags
    cleanr = re.compile('<.*?>$< /><')
    #cleanr = re.compile('<.*?>')
    cleantext = re.sub(cleanr, ' ', sentence)
    return cleantext
def cleanpunc(sentence): #function to clean the word of any punctuation
or special characters
    cleaned = re.sub(r'[?|!|\\'|"|#]',r'',sentence)
    cleaned = re.sub(r'[,|,)|(|\\|/]',r' ',cleaned)
    return cleaned

```

In [0]: #Code for implementing step-by-step the checks mentioned in the pre-pro
cessing phase.

```

'''Pre processing of text data:It is cleaning and flitering text'''
i=0
str1=' '
global final_string
final_string=[]
all_positive_words=[]
all_negative_words=[]
s=''
for sent in final['Text'].values:
    filtered_sentence=[]
    #print(sent);
    sent=cleanhtml(sent) # remove HTML tags
    for w in sent.split():
        for cleaned_words in cleanpunc(w).split():
            if((cleaned_words.isalpha()) & (len(cleaned_words)>2)):
                if(cleaned_words.lower() not in stop):
                    s=(sno.stem(cleaned_words.lower())).encode('utf8')
                    filtered_sentence.append(s)
                    if (final['Score'].values)[i] == 1:
                        all_positive_words.append(s) #list of all words
used to describe positive reviews

```

```

        if(final['Score'].values)[i] == 0:
            all_negative_words.append(s) #list of all words
used to describe negative reviews reviews
        else:
            continue
    else:
        continue
    #print(filtered_sentence)
    str1 = b" ".join(filtered_sentence) #final string of cleaned words
    #print("*****")
    *****

    final_string.append(str1)
    i+=1

```

Dumping and loading Pre processing of text data in pickle file

```

In [0]: pickle_path_final_string='final_string.pkl'
        final_string_file=open(pickle_path_final_string,'wb')
        pickle.dump(final_string,final_string_file)
        final_string_file.close()

```

```

In [0]: pickle_path_final_string='final_string.pkl'
        final_string_unpkl=open(pickle_path_final_string,'rb')
        final_string=pickle.load(final_string_unpkl)

```

```

In [0]: final['CleanedText']=final_string
        final['CleanedText']=final['CleanedText'].str.decode("utf-8")
        #below the processed review can be seen in the CleanedText Column
        print('Shape of final',final.shape)
        final.head()

```

Shape of final (393761, 11)

Out[0]:

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	He
--	----	-----------	--------	-------------	----------------------	----

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	He
150528	150529	0006641040	A25ACLV5KPB4W	Matt Hetling "Matt"	0	1
150506	150507	0006641040	A1S4A3IQ2MU7V4	sally sue "sally sue"	1	1
150505	150506	0006641040	A2IW4PEEKO2R0U	Tracy	1	1
150504	150505	0006641040	A2PTSM496CF40Z	Jason A. Teeple "Nobody made a greater mistak...	1	1
150503	150504	0006641040	AQEYF1AXARWJZ	Les Sinclair "book maven"	1	1

Converting Reviews Data into IMDB Dataset

```
In [0]: ##Sorting data according to Time in ascending order
## Time Based Splitting
time_sorted_data = final.sort_values('Time', axis=0,
                                     ascending=True,
                                     inplace=False,
                                     kind='quicksort',
                                     na_position='last')

x = time_sorted_data['CleanedText'].values
y = time_sorted_data['Score']

# vocabulary of words using countvectorizer
count_vect = CountVectorizer()
count_vect.fit(x)

vocabulary = count_vect.get_feature_names()
print('No. of words in the Vocabulary : ', len(vocabulary))
```

No. of words in the Vocabulary : 73405

```
In [0]: # Storing words in vocabulary
word_dict=dict()
index=0
for words in x:
    for word in words.split():
        word_dict.setdefault(word, [])
        word_dict[word].append(index)
        index+=1

# Getting frequency for each word of vocabulary

freq = []
for w in vocabulary:
    freq.append(len(word_dict[w]))
```

```
print("done")
```

done

```
In [0]: # sort the frequency in decending order
```

```
new_index=np.argsort(np.array(freq))[::-1]
```

```
In [0]: # Allocate rank to freq in decending order
```

```
word_rank=dict()
```

```
rank=1
```

```
for i in new_index:
```

```
    word_rank[vocabulary[i]]=rank
```

```
    rank+=1
```

converting data into imdb format

```
In [0]: review_data=[]
```

```
for sent in x:
```

```
    row=[]
```

```
    for word in sent.split():
```

```
        if len(word)>1 :
```

```
            row.append(word_rank[word])
```

```
    review_data.append(row)
```

```
In [0]: # Splitting the data into 60-40 train_data and test_data
```

```
from sklearn.model_selection import train_test_split
```

```
X_train, X_test, Y_train, Y_test = train_test_split(review_data, y,  
                                                    test_size=0.5,  
                                                    random_state=42)
```

```
In [0]: print("length of X_train :",len(X_train))
```

```
print("length of X_test :",len(X_test))
```

```
print("Shape of Y_train :",Y_train.shape)
```

```
print("Shape of Y_test :",Y_test.shape)
```

```
length of X_train : 196880
length of X_test : 196881
Shape of Y_train : (196880,)
Shape of Y_test : (196881,)
```

```
In [0]: print( review_data[0])
```

```
[13351, 27, 1106, 14, 379, 2444, 3182, 13972, 1142, 1241, 545, 109, 384
5, 7102, 860, 800, 10363, 1852, 11939, 8, 174, 780, 1106, 1082, 2807, 1
211, 1106, 1429, 1860, 379, 79, 263, 13972, 1438, 1699]
```

```
In [0]: # truncate and/or pad input sequences
```

```
max_review_length = 120
X_train = sequence.pad_sequences(X_train, maxlen=max_review_length)
X_test = sequence.pad_sequences(X_test, maxlen=max_review_length)
```

```
print(X_train.shape)
print(X_train[1])
print(X_test.shape)
```

```
(196880, 120)
[  0  0  0  0  0  0  0  0  0  0  0  0
0  0  0  0  0  0  0  0  0  0  0  0
0  0  0  0  0  0  0  0  0  0  0  0
0  0  0  0  0  0  0  0  0  0  0  0
0  0  0  0  0  0  0  0  0  0  0  0
0  0  0  0  0  0  0  0  0  0  0  0
0  0  0  0  0  0  0  0  0  0  0  0
0  0  0  0  0  0  0  0  0  0  0  0
7 293 1255 2912 478 29 41 26 639 2978 412 3887 91
9 1305 1753 87 287 378 328 582 172 419 326 117 4
1
```

```
2610 410 1399 826 631 113 1386 1014 10807 478 62 5
0]
(196881, 120)
```

```
In [0]: len_voc=len(vocabulary)+1
        print(len_voc)
```

```
73406
```

1 RNN with 1 LSTM layer

```
In [0]: start = datetime.now()
        embedding_len=32
        lstm_1=Sequential()
        lstm_1.add(Embedding(len_voc,
                               embedding_len,
                               input_length=max_review_length))

        lstm_1.add(Dropout(0.15))

        lstm_1.add(LSTM(100))# Adding first LSTM layer

        lstm_1.add(Dropout(0.25))# Adding Dropout

        lstm_1.add(Dense(1, activation='sigmoid'))# Adding output layer

        # Printing the model summary
        print(lstm_1.summary())

        # Compiling the model
        lstm_1.compile(loss='binary_crossentropy',
                        optimizer='adam',
                        metrics=['accuracy'])

        # Fitting the data to the model
```

```

history_1 = lstm_1.fit(X_train, Y_train, nb_epoch=10,
                      batch_size=512 ,verbose=1,
                      validation_data=(X_test, Y_test))

#evaluating model
score=lstm_1.evaluate(X_test,Y_test,verbose=0)
test_score_lstm_1=score[0]
test_accuracy_lstm_1=score[1]
train_accuracy_lstm_1=max(history_1.history['acc'])
print('test score :',test_score_lstm_1)
print('test Accuracy : ',test_accuracy_lstm_1)

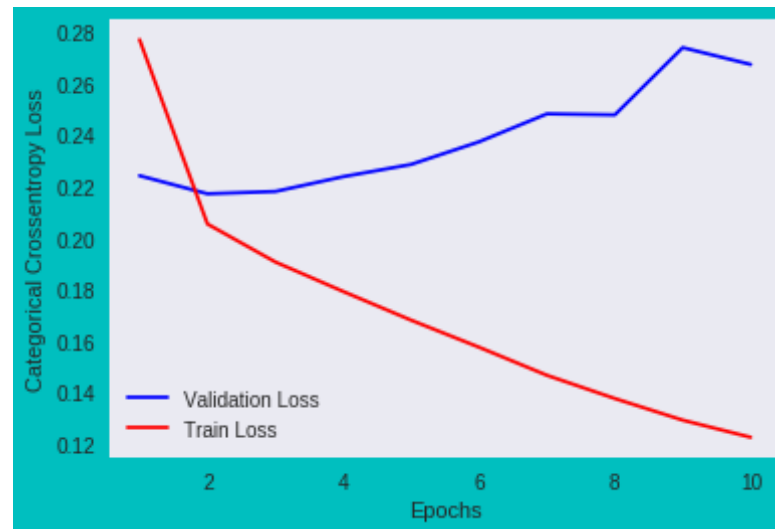
# error plot
x=list(range(1,11))
vy=history_1.history['val_loss'] #validation loss
ty=history_1.history['loss'] # train loss
plt_dynamic(x, vy, ty)

runtime_lstm1=datetime.now() - start
print("Time taken to run this cell :", runtime_lstm1)

```

Layer (type)	Output Shape	Param #
embedding_2 (Embedding)	(None, 120, 32)	2348992
dropout_2 (Dropout)	(None, 120, 32)	0
lstm_2 (LSTM)	(None, 100)	53200
dropout_3 (Dropout)	(None, 100)	0
dense_2 (Dense)	(None, 1)	101
Total params: 2,402,293		
Trainable params: 2,402,293		
Non-trainable params: 0		
None		
Train on 196880 samples, validate on 196881 samples		
Epoch 1/10		

```
196880/196880 [=====] - 518s 3ms/step - loss:
0.2768 - acc: 0.8921 - val_loss: 0.2240 - val_acc: 0.9087
Epoch 2/10
196880/196880 [=====] - 514s 3ms/step - loss:
0.2053 - acc: 0.9167 - val_loss: 0.2170 - val_acc: 0.9118
Epoch 3/10
196880/196880 [=====] - 514s 3ms/step - loss:
0.1906 - acc: 0.9228 - val_loss: 0.2179 - val_acc: 0.9115
Epoch 4/10
196880/196880 [=====] - 514s 3ms/step - loss:
0.1791 - acc: 0.9273 - val_loss: 0.2237 - val_acc: 0.9081
Epoch 5/10
196880/196880 [=====] - 513s 3ms/step - loss:
0.1680 - acc: 0.9320 - val_loss: 0.2284 - val_acc: 0.9097
Epoch 6/10
196880/196880 [=====] - 514s 3ms/step - loss:
0.1575 - acc: 0.9369 - val_loss: 0.2372 - val_acc: 0.9095
Epoch 7/10
196880/196880 [=====] - 515s 3ms/step - loss:
0.1467 - acc: 0.9412 - val_loss: 0.2480 - val_acc: 0.9067
Epoch 8/10
196880/196880 [=====] - 513s 3ms/step - loss:
0.1377 - acc: 0.9453 - val_loss: 0.2476 - val_acc: 0.9046
Epoch 9/10
196880/196880 [=====] - 513s 3ms/step - loss:
0.1293 - acc: 0.9487 - val_loss: 0.2737 - val_acc: 0.9039
Epoch 10/10
196880/196880 [=====] - 513s 3ms/step - loss:
0.1226 - acc: 0.9516 - val_loss: 0.2672 - val_acc: 0.9036
test score : 0.26717215862544025
test Accuracy : 0.9035559551207034
```



Time taken to run this cell : 1:27:53.818856

Observation:

- After looking at training and testing error plot, model can be diagnosed as overfitting.
- Training error slopes down while validation error also slopes and hit at inflection point.
- After the inflection point, validation loss slopes continuously increasing.

2 RNN with 2 LSTM layer

```
In [0]: from keras.layers.normalization import BatchNormalization
```

```
In [0]: start = datetime.now()
embedding_len=32
lstm_2=Sequential()
lstm_2.add(Embedding(len_voc,
                    embedding_len,
                    input_length=max_review_length))
```



```

lstm_2.add(Dropout(0.15))

lstm_2.add(LSTM(100,return_sequences=True,
               dropout=0.15,
               recurrent_dropout=0.25))# Adding first LSTM layer

lstm_2.add(BatchNormalization())
lstm_2.add(LSTM(100,
               dropout=0.15,
               recurrent_dropout=0.25))# Adding second LSTM layer


#lstm_2.add(BatchNormalization())
lstm_2.add(Dense(1, activation='sigmoid'))# Adding output layer

# Printing the model summary
print(lstm_2.summary())

# Compiling the model
lstm_2.compile(loss='binary_crossentropy',
               optimizer='adam',
               metrics=['accuracy'])

# Fitting the data to the model
history_2 = lstm_2.fit(X_train, Y_train, nb_epoch=10,
                      batch_size=512 ,verbose=1,
                      validation_data=(X_test, Y_test))
runtime_lstm2=datetime.now() - start
print("Time taken to run this cell :", runtime_lstm2)

```

Layer (type)	Output Shape	Param #
embedding_1 (Embedding)	(None, 120, 32)	2348992
dropout_1 (Dropout)	(None, 120, 32)	0
lstm_1 (LSTM)	(None, 120, 100)	53200

batch_normalization_1 (Batch Normalization)	(None, 120, 100)	400
lstm_2 (LSTM)	(None, 100)	80400
dense_1 (Dense)	(None, 1)	101


Total params: 2,483,093
 Trainable params: 2,482,893
 Non-trainable params: 200

None
 Train on 196880 samples, validate on 196881 samples
 Epoch 1/10
 196880/196880 [=====] - 1611s 8ms/step - loss: 0.2601 - acc: 0.8945 - val_loss: 0.2188 - val_acc: 0.9100
 Epoch 2/10
 196880/196880 [=====] - 1596s 8ms/step - loss: 0.2043 - acc: 0.9166 - val_loss: 0.2135 - val_acc: 0.9126
 Epoch 3/10
 196880/196880 [=====] - 1598s 8ms/step - loss: 0.1856 - acc: 0.9242 - val_loss: 0.2161 - val_acc: 0.9115
 Epoch 4/10
 196880/196880 [=====] - 1599s 8ms/step - loss: 0.1699 - acc: 0.9311 - val_loss: 0.2251 - val_acc: 0.9111
 Epoch 5/10
 196880/196880 [=====] - 1596s 8ms/step - loss: 0.1566 - acc: 0.9369 - val_loss: 0.2413 - val_acc: 0.9114
 Epoch 6/10
 196880/196880 [=====] - 1600s 8ms/step - loss: 0.1450 - acc: 0.9415 - val_loss: 0.2553 - val_acc: 0.9096
 Epoch 7/10
 196880/196880 [=====] - 1596s 8ms/step - loss: 0.1350 - acc: 0.9458 - val_loss: 0.2429 - val_acc: 0.9088
 Epoch 8/10
 196880/196880 [=====] - 1598s 8ms/step - loss: 0.1261 - acc: 0.9495 - val_loss: 0.2651 - val_acc: 0.9091
 Epoch 9/10
 196880/196880 [=====] - 1600s 8ms/step - loss:

```
0.1187 - acc: 0.9524 - val_loss: 0.2743 - val_acc: 0.9037
Epoch 10/10
196880/196880 [=====] - 1599s 8ms/step - loss:
0.1116 - acc: 0.9550 - val_loss: 0.3064 - val_acc: 0.9087
Time taken to run this cell : 4:26:38.657956
```

```
In [0]: #evaluating model
score=lstm_2.evaluate(X_test,Y_test,verbose=0)
test_score_lstm_2=score[0]
test_accuracy_lstm_2=score[1]
train_accuracy_lstm_2=max(history_2.history['acc'])
print('test score :',test_score_lstm_2)
print('test Accuracy : ',test_accuracy_lstm_2)
# error plot
x=list(range(1,11))
vy=history_2.history['val_loss'] #validation loss
ty=history_2.history['loss'] # train loss
plt_dynamic(x, vy, ty)
```

```
test score : 0.3063991781577867
test Accuracy : 0.9087469080313448
```



3 RNN with 4 LSTM layer

```
In [0]: start = datetime.now()
embedding_len=32
lstm_4=Sequential()
lstm_4.add(Embedding(len_voc,
                    embedding_len,
                    input_length=max_review_length))

lstm_4.add(Dropout(0.25))

lstm_4.add(LSTM(100,return_sequences=True,
              dropout=0.15,
              recurrent_dropout=0.25))# Adding first LSTM layer
lstm_4.add(LSTM(80,return_sequences=True,
              dropout=0.15,
              recurrent_dropout=0.25))# Adding second LSTM layer

lstm_4.add(BatchNormalization())
lstm_4.add(LSTM(50,return_sequences=True,
              dropout=0.15,
              recurrent_dropout=0.25))# Adding Third LSTM layer
lstm_4.add(BatchNormalization())
lstm_4.add(LSTM(40,
              dropout=0.15))# Adding Fourth LSTM layer
lstm_4.add(BatchNormalization())

lstm_4.add(Dense(1, activation='sigmoid'))# Adding output layer

# Printing the model summary
print(lstm_4.summary())

# Compiling the model
lstm_4.compile(loss='binary_crossentropy',
              optimizer='adam',
```

```

metrics=['accuracy'])

# Fitting the data to the model
history_4 = lstm_4.fit(X_train, Y_train, nb_epoch=10,
                      batch_size=512 ,verbose=1,
                      validation_data=(X_test, Y_test))
runtime_lstm4=datetime.now() - start
print("Time taken to run this cell :", runtime_lstm4)

```

Layer (type)	Output Shape	Param #
embedding_1 (Embedding)	(None, 120, 32)	2348992
dropout_1 (Dropout)	(None, 120, 32)	0
lstm_1 (LSTM)	(None, 120, 100)	53200
lstm_2 (LSTM)	(None, 120, 80)	57920
batch_normalization_1 (Batch Normalization)	(None, 120, 80)	320
lstm_3 (LSTM)	(None, 120, 50)	26200
batch_normalization_2 (Batch Normalization)	(None, 120, 50)	200
lstm_4 (LSTM)	(None, 40)	14560
batch_normalization_3 (Batch Normalization)	(None, 40)	160
dense_1 (Dense)	(None, 1)	41

=====
 Total params: 2,501,593
 Trainable params: 2,501,253
 Non-trainable params: 340
 =====

None
 Train on 196880 samples, validate on 196881 samples
 Epoch 1/10
 196880/196880 [=====] - 2089s 11ms/step - los

```

s: 0.3766 - acc: 0.8375 - val_loss: 0.2369 - val_acc: 0.9066
Epoch 2/10
196880/196880 [=====] - 2081s 11ms/step - los
s: 0.2255 - acc: 0.9095 - val_loss: 0.2283 - val_acc: 0.9058
Epoch 3/10
196880/196880 [=====] - 2087s 11ms/step - los
s: 0.2071 - acc: 0.9165 - val_loss: 0.2195 - val_acc: 0.9096
Epoch 4/10
196880/196880 [=====] - 2091s 11ms/step - los
s: 0.1967 - acc: 0.9209 - val_loss: 0.2226 - val_acc: 0.9094
Epoch 5/10
196880/196880 [=====] - 2090s 11ms/step - los
s: 0.1843 - acc: 0.9261 - val_loss: 0.2223 - val_acc: 0.9107
Epoch 6/10
196880/196880 [=====] - 2085s 11ms/step - los
s: 0.1757 - acc: 0.9296 - val_loss: 0.2269 - val_acc: 0.9115
Epoch 7/10
196880/196880 [=====] - 2080s 11ms/step - los
s: 0.1677 - acc: 0.9328 - val_loss: 0.2349 - val_acc: 0.9041
Epoch 8/10
196880/196880 [=====] - 2081s 11ms/step - los
s: 0.1600 - acc: 0.9363 - val_loss: 0.2379 - val_acc: 0.9096
Epoch 9/10
196880/196880 [=====] - 2082s 11ms/step - los
s: 0.1525 - acc: 0.9392 - val_loss: 0.2378 - val_acc: 0.9086
Epoch 10/10
196880/196880 [=====] - 2084s 11ms/step - los
s: 0.1461 - acc: 0.9424 - val_loss: 0.2522 - val_acc: 0.9075
Time taken to run this cell : 5:47:39.955757

```

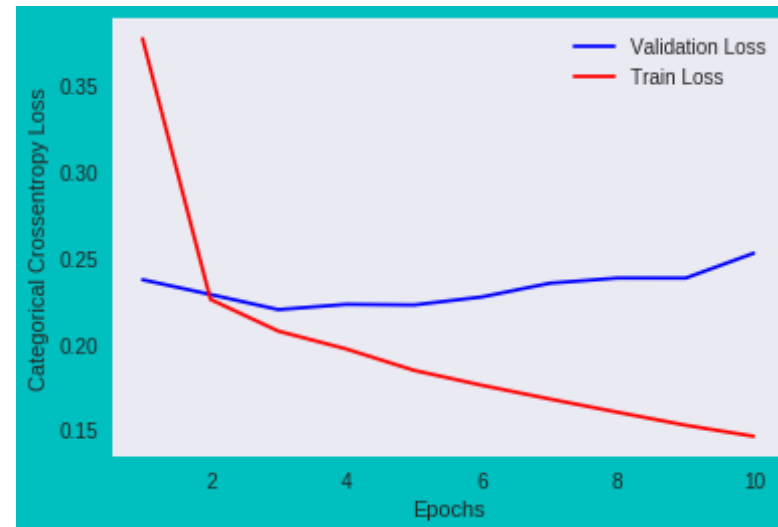
```

In [0]: #evaluating model
score=lstm_4.evaluate(X_test,Y_test,verbose=0)
test_score_lstm_4=score[0]
test_accuracy_lstm_4=score[1]
train_accuracy_lstm_4=max(history_4.history['acc'])
print('test score :',test_score_lstm_4)
print('test Accuracy :',test_accuracy_lstm_4)
# error plot
x=list(range(1,11))

```

```
vy=history_4.history['val_loss'] #validation loss
ty=history_4.history['loss'] # train loss
plt_dynamic(x, vy, ty)
```

test score : 0.2522423982056529
test Accuracy : 0.9075177391425236



4 RNN with 5 LSTM layer

```
In [0]: start = datetime.now()
embedding_len=32
lstm_5=Sequential()
lstm_5.add(Embedding(len_voc,
                    embedding_len,
                    input_length=max_review_length))

lstm_5.add(Dropout(0.45))

lstm_5.add(LSTM(120,return_sequences=True,
               dropout=0.35,
               recurrent_dropout=0.25))# Adding first LSTM layer
```

```

lstm_5.add(LSTM(100,return_sequences=True,
               dropout=0.35,
               recurrent_dropout=0.35))# Adding second LSTM layer
lstm_5.add(BatchNormalization())
lstm_5.add(LSTM(100,return_sequences=True,
               dropout=0.35,
               recurrent_dropout=0.35))# Adding Third LSTM layer
lstm_5.add(BatchNormalization())
lstm_5.add(LSTM(80,return_sequences=True,
               dropout=0.35,
               recurrent_dropout=0.35))# Adding Fourth LSTM layer
lstm_5.add(BatchNormalization())
lstm_5.add(LSTM(60,dropout=0.35,
               recurrent_dropout=0.35))# Adding Fifth LSTM layer
lstm_5.add(BatchNormalization())
lstm_5.add(Dense(1, activation='sigmoid'))# Adding output layer

# Printing the model summary
print(lstm_5.summary())

# Compiling the model
lstm_5.compile(loss='binary_crossentropy',
               optimizer='adam',
               metrics=['accuracy'])

# Fitting the data to the model
history_5 = lstm_5.fit(X_train, Y_train, nb_epoch=10,
                      batch_size=512 ,verbose=1,
                      validation_data=(X_test, Y_test))

#evaluating model
score=lstm_5.evaluate(X_test,Y_test,verbose=0)
test_score_lstm_5=score[0]
test_accuracy_lstm_5=score[1]
train_accuracy_lstm_5=max(history_5.history['acc'])
print("Train Accuracy",train_accuracy_lstm_5)
print('test score :',test_score_lstm_5)
print('test Accuracy :',test_accuracy_lstm_5)
# error plot
x=list(range(1,11))

```



```
vy=history_5.history['val_loss'] #validation loss
ty=history_5.history['loss'] # train loss
plt_dynamic(x, vy, ty)
runtime_lstm5=datetime.now() - start
print("Time taken to run this cell :", runtime_lstm5)
```

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow/python/framework/op_def_library.py:263: colocate_with (from tensorflow.python.framework.ops) is deprecated and will be removed in a future version.

Instructions for updating:

Colocations handled automatically by placer.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:3445: calling dropout (from tensorflow.python.ops.nn_ops) with keep_prob is deprecated and will be removed in a future version.

Instructions for updating:

Please use `rate` instead of `keep_prob`. Rate should be set to `rate = 1 - keep_prob`.

Layer (type)	Output Shape	Param #
embedding_1 (Embedding)	(None, 120, 32)	2348992
dropout_1 (Dropout)	(None, 120, 32)	0
lstm_1 (LSTM)	(None, 120, 120)	73440
lstm_2 (LSTM)	(None, 120, 100)	88400
batch_normalization_1 (Batch Normalization)	(None, 120, 100)	400
lstm_3 (LSTM)	(None, 120, 100)	80400
batch_normalization_2 (Batch Normalization)	(None, 120, 100)	400
lstm_4 (LSTM)	(None, 120, 80)	57920
batch_normalization_3 (Batch Normalization)	(None, 120, 80)	320

lstm_5 (LSTM)	(None, 60)	33840
<hr/>		
batch_normalization_4 (Batch Normalization)	(None, 60)	240
<hr/>		
dense_1 (Dense)	(None, 1)	61
<hr/>		
Total params: 2,684,413		
Trainable params: 2,683,733		
Non-trainable params: 680		

None

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow/python/ops/math_ops.py:3066: to_int32 (from tensorflow.python.ops.math_ops) is deprecated and will be removed in a future version.

Instructions for updating:

Use tf.cast instead.

Train on 196880 samples, validate on 196881 samples

Epoch 1/10

196880/196880 [=====] - 3164s 16ms/step - loss: 0.4910 - acc: 0.8030 - val_loss: 0.4144 - val_acc: 0.8547

Epoch 2/10

196880/196880 [=====] - 3130s 16ms/step - loss: 0.4248 - acc: 0.8552 - val_loss: 0.4166 - val_acc: 0.8547

Epoch 3/10

196880/196880 [=====] - 3172s 16ms/step - loss: 0.4192 - acc: 0.8553 - val_loss: 0.4237 - val_acc: 0.8547

Epoch 4/10

196880/196880 [=====] - 3161s 16ms/step - loss: 0.3334 - acc: 0.8690 - val_loss: 0.2458 - val_acc: 0.8992

Epoch 5/10

196880/196880 [=====] - 3140s 16ms/step - loss: 0.2501 - acc: 0.8991 - val_loss: 0.2314 - val_acc: 0.9070

Epoch 6/10

196880/196880 [=====] - 3148s 16ms/step - loss: 0.2285 - acc: 0.9081 - val_loss: 0.2320 - val_acc: 0.9095

Epoch 7/10


196880/196880 [=====] - 3145s 16ms/step - loss: 0.2169 - acc: 0.9122 - val_loss: 0.2226 - val_acc: 0.9090

Epoch 8/10

```

196880/196880 [=====] - 3144s 16ms/step - los
s: 0.2092 - acc: 0.9161 - val_loss: 0.2308 - val_acc: 0.9100
Epoch 9/10
196880/196880 [=====] - 3182s 16ms/step - los
s: 0.2053 - acc: 0.9176 - val_loss: 0.2244 - val_acc: 0.9104
Epoch 10/10
196880/196880 [=====] - 3178s 16ms/step - los
s: 0.1990 - acc: 0.9202 - val_loss: 0.2235 - val_acc: 0.9110
Train Accuracy 0.9202356765784659
test score : 0.2235467241037283
test Accuracy : 0.9109716021359054

```



Time taken to run this cell : 9:02:45.687949

```
In [0]: print('Train loss',ty)
```

```

Train loss [0.49096747402106294, 0.4247921708370311, 0.419169231686152
3, 0.33343958246519, 0.2501070746763699, 0.2285251474087435, 0.21694653
656777207, 0.20923510087015956, 0.20528901007566835, 0.1989786412913544
5]

```

```
In [0]: print('test score',test_score_lstm_5)
```

```
test score : 0.2235467241037283
```

test score 0.223546/24103/283

Observation

```
In [0]: from prettytable import PrettyTable
models=['RNN With 1 LSTM Layer',
         'RNN With 2 LSTM Layer',
         'RNN With 4 LSTM Layer',
         'RNN With 5 LSTM Layer']
training_accuracy=[train_accuracy_lstm_1,train_accuracy_lstm_2,
                   train_accuracy_lstm_4,train_accuracy_lstm_5]
test_accuracy=[test_accuracy_lstm_1,test_accuracy_lstm_2,
               test_accuracy_lstm_4,test_accuracy_lstm_5]
INDEX = [1,2,3,4]
# Initializing prettytable
Model_Performance = PrettyTable()
# Adding columns
Model_Performance.add_column("INDEX.",INDEX)
Model_Performance.add_column("MODEL_NAME",models)
Model_Performance.add_column("TRAINING ACCURACY",training_accuracy)
Model_Performance.add_column("TESTING ACCURACY",test_accuracy)
#Model_Performance.add_column("TEST SCORE",test_score)

# Printing the Model_Performance
print(Model_Performance)
```

```
+-----+-----+-----+-----+
-+
| INDEX. |      MODEL_NAME      | TRAINING ACCURACY | TESTING ACCURACY |
|
+-----+-----+-----+-----+
-+
| 1      | RNN With 1 LSTM Layer |      0.9423      |      0.9035      |
|
| 2      | RNN With 2 LSTM Layer |      0.9392      |      0.9087      |
|
| 3      | RNN With 4 LSTM Layer |    0.936954      |      0.9075      |
|
```

	4	RNN With 5 LSTM Layer	0.9202		0.9108
+	-	-	-	-	-
-	+				